

## Session 2: Assignment 1

### **Task 1:**

We have a dataset of sales of different TV sets across different locations.

The fields are arranged like: Company Name|Product Name|Size in inches|State|Pin Code|Price  
Records look like: Samsung|Optima|14|Madhya Pradesh|132401|14200

There are some invalid records which contain 'NA' in either Company Name or Product Name.

**1. Write a MapReduce program to filter out the invalid records. Map only job will fit for this context.**

Output for filter out invalid records: Total number of input count was 18 which consist of 2 invalid records. After removing the count of output is 16. Below the screenshot of the output.

```
[acadgild@localhost ~]$ hadoop jar tv set.jar com.tv set.tv set /tv set.txt /out tv8
18/07/29 21:42:59 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where
18/07/29 21:43:01 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/29 21:43:02 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and
his.
18/07/29 21:43:02 INFO input.FileInputFormat: Total input paths to process : 1
18/07/29 21:43:03 INFO mapreduce.JobSubmitter: number of splits:1
18/07/29 21:43:03 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532876630081_0009
18/07/29 21:43:03 INFO impl.YarnClientImpl: Submitted application application_1532876630081_0009
18/07/29 21:43:03 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532876630081_0009/
18/07/29 21:43:03 INFO mapreduce.Job: Running job: job_1532876630081_0009
18/07/29 21:43:14 INFO mapreduce.Job: Job job_1532876630081_0009 running in uber mode : false
18/07/29 21:43:14 INFO mapreduce.Job:  map 0% reduce 0%
18/07/29 21:43:21 INFO mapreduce.Job:  map 100% reduce 0%
18/07/29 21:43:29 INFO mapreduce.Job:  map 100% reduce 100%
18/07/29 21:43:29 INFO mapreduce.Job: Job job_1532876630081_0009 completed successfully
18/07/29 21:43:29 INFO mapreduce.Job: Counters: 49
    File System Counters
      FILE: Number of bytes read=204
      FILE: Number of bytes written=216043
      FILE: Number of read operations=0
      FILE: Number of large read operations=0
      FILE: Number of write operations=0
      HDFS: Number of bytes read=814
      HDFS: Number of bytes written=150
      HDFS: Number of read operations=6
      HDFS: Number of large read operations=0
      HDFS: Number of write operations=2
    Job Counters
```

```

WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=717
File Output Format Counters
  Bytes Written=150
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /out tv8/part-r-00000
18/07/29 21:43:38 WARN util.NativeCodeLoader: Unable to load native-hadoop library
Akai      16
Lava      20
Lava      20
Lava      20
Onida     14
Onida     18
Onida     18
Samsung   14
Samsung   14
Samsung   14
Samsung   16
Samsung   14
Samsung   14
Samsung   14
Samsung   14
Zen       14
Zen       14
[acadgild@localhost ~]$ █

```

## 2. Write a MapReduce program to calculate the total units sold for each Company

### Output for total units sold for each company

```

[acadgild@localhost ~]$ hadoop jar tv set.jar com.tv set.tv set /tv set.txt /out tv9
18/07/29 21:57:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java class
18/07/29 21:57:48 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/29 21:57:49 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool inter
nis.
18/07/29 21:57:50 INFO input.FileInputFormat: Total input paths to process : 1
18/07/29 21:57:50 INFO mapreduce.JobSubmitter: number of splits:1
18/07/29 21:57:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532876630081_0010
18/07/29 21:57:50 INFO impl.YarnClientImpl: Submitted application application_1532876630081_0010
18/07/29 21:57:51 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532876630081_0010/
18/07/29 21:57:51 INFO mapreduce.Job: Running job: job_1532876630081_0010
18/07/29 21:58:00 INFO mapreduce.Job: Job job_1532876630081_0010 running in uber mode : false
18/07/29 21:58:00 INFO mapreduce.Job:  map 0% reduce 0%
18/07/29 21:58:07 INFO mapreduce.Job:  map 100% reduce 0%
18/07/29 21:58:15 INFO mapreduce.Job:  map 100% reduce 100%
18/07/29 21:58:16 INFO mapreduce.Job: Job job_1532876630081_0010 completed successfully
18/07/29 21:58:16 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=204
  FILE: Number of bytes written=216363
  FILE: Number of read operations=0

```

```

        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=717
    File Output Format Counters
        Bytes Written=44
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /out tv9/part-r-00000
18/07/29 21:58:26 WARN util.NativeCodeLoader: Unable to load native-hadoop library
Akai      16
Lava      60
Onida     50
Samsung   100
Zen       28
[acadgild@localhost ~]$ █

```



**3. Write a MapReduce program to calculate the total units sold in each state for Onida company.**

**Output to calculate total units sold by Onida company in each state.**

```

[acadgild@localhost ~]$ hadoop jar saletv.jar com.sale.sale tv /tv set.txt /sale tv out
18/07/29 23:04:51 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
18/07/29 23:04:53 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/29 23:04:54 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool in
his.
18/07/29 23:04:55 INFO input.FileInputFormat: Total input paths to process : 1
18/07/29 23:04:55 INFO mapreduce.JobSubmitter: number of splits:1
18/07/29 23:04:55 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532876630081_0011
18/07/29 23:04:56 INFO impl.YarnClientImpl: Submitted application application_1532876630081_0011
18/07/29 23:04:56 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532876630081_0011/
18/07/29 23:04:56 INFO mapreduce.Job: Running job: job_1532876630081_0011
18/07/29 23:05:07 INFO mapreduce.Job: Job job_1532876630081_0011 running in uber mode : false
18/07/29 23:05:07 INFO mapreduce.Job: map 0% reduce 0%
18/07/29 23:05:14 INFO mapreduce.Job: map 100% reduce 0%
18/07/29 23:05:22 INFO mapreduce.Job: map 100% reduce 100%
18/07/29 23:05:23 INFO mapreduce.Job: Job job_1532876630081_0011 completed successfully
18/07/29 23:05:24 INFO mapreduce.Job: Counters: 49

```

```

Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=717
File Output Format Counters
    Bytes Written=17
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /sale tv out/part-r-00000
18/07/29 23:05:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform.
Uttar Pradesh 50

```

## Task 2:

Dataset is sample data of songs heard by users on an online streaming platform. The Description of data set attached in musicdata.txt is as follows: - Col\_1 - UserId Col\_2 - TrackId Col\_3 - Songs Share status (1 for shared, 0 for not shared) Col\_4 - Listening Platform (Radio or Web - 0 for radio, 1 for web) Col\_5 - Song Listening Status (0 for skipped, 1 for fully heard) 1. Write MapReduce program for following tasks.

**Task 1 Find the number of unique listeners in the data set.**

**Output for Unique Listener**

```

[acadgild@localhost ~]$ hadoop jar radio23.jar com.mapreduce.radio /music6.txt /temp out19
18/07/28 20:39:21 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/07/28 20:39:22 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/28 20:39:23 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute this.
18/07/28 20:39:24 INFO input.FileInputFormat: Total input paths to process : 1
18/07/28 20:39:24 INFO mapreduce.JobSubmitter: number of splits:1
18/07/28 20:39:24 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532787160984_0014
18/07/28 20:39:24 INFO impl.YarnClientImpl: Submitted application application_1532787160984_0014
18/07/28 20:39:24 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532787160984_0014/
18/07/28 20:39:24 INFO mapreduce.Job: Running job: job_1532787160984_0014
18/07/28 20:39:36 INFO mapreduce.Job: Job job_1532787160984_0014 running in uber mode : false
18/07/28 20:39:36 INFO mapreduce.Job: map 0% reduce 0%
18/07/28 20:39:43 INFO mapreduce.Job: map 100% reduce 0%
18/07/28 20:39:51 INFO mapreduce.Job: map 100% reduce 100%
18/07/28 20:39:52 INFO mapreduce.Job: Job job_1532787160984_0014 completed successfully
18/07/28 20:39:52 INFO mapreduce.Job: Counters: 49
File System Counters
    FILE: Number of bytes read=46
    FILE: Number of bytes written=216087
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=165

```

```

        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=68
    File Output Format Counters
        Bytes Written=27
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -ls /temp_out19
18/07/28 20:40:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
Found 2 items
-rw-r--r--   1 acadgild supergroup          0 2018-07-28 20:39 /temp_out19/_SUCCESS
-rw-r--r--   1 acadgild supergroup        27 2018-07-28 20:39 /temp_out19/part-r-00000
[acadgild@localhost ~]$ hadoop fs -cat /temp_out19/part-r-00000
18/07/28 20:40:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
111113 1
111115 2
111117 1
[acadgild@localhost ~]$

```

**Task 2 What are the number of times a song was heard fully.**

**Output for song listened fully**

```

223 1
225 2
[acadgild@localhost ~]$ hadoop jar song.jar com.songListen.songListen /music6.txt /song output
18/07/29 01:40:28 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
18/07/29 01:40:29 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/29 01:40:30 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool
his.
18/07/29 01:40:30 INFO input.FileInputFormat: Total input paths to process : 1
18/07/29 01:40:30 INFO mapreduce.JobSubmitter: number of splits:1
18/07/29 01:40:31 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532787160984_0035
18/07/29 01:40:31 INFO impl.YarnClientImpl: Submitted application application_1532787160984_0035
18/07/29 01:40:31 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532787160984_0035/
18/07/29 01:40:31 INFO mapreduce.Job: Running job: job_1532787160984_0035
18/07/29 01:40:42 INFO mapreduce.Job: Job job_1532787160984_0035 running in uber mode : false
18/07/29 01:40:42 INFO mapreduce.Job: map 0% reduce 0%
18/07/29 01:40:50 INFO mapreduce.Job: map 100% reduce 0%
18/07/29 01:40:59 INFO mapreduce.Job: map 100% reduce 100%
18/07/29 01:40:59 INFO mapreduce.Job: Job job_1532787160984_0035 completed successfully
18/07/29 01:40:59 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=16
        FILE: Number of bytes written=216063
        ...
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=68
    File Output Format Counters
        Bytes Written=6
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /song output/part-r-00000
18/07/29 01:41:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
223 1
[acadgild@localhost ~]$ hadoop fs -cat /music6.txt
18/07/29 01:42:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
111115|222|0|1|0
111113|225|1|0|0
111117|223|0|1|1
111115|225|1|0|0
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

### TASK 3 What are the number of times a song was shared.

#### Output for song shared

```
acadgild@localhost:~$  
You have new mail in /var/spool/mail/acadgild  
[acadgild@localhost ~]$ hadoop jar songshared.jar songShared /music6.txt /songShared output  
18/07/29 02:05:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j  
18/07/29 02:05:35 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032  
18/07/29 02:05:36 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the To  
his.  
18/07/29 02:05:37 INFO input.FileInputFormat: Total input paths to process : 1  
18/07/29 02:05:37 INFO mapreduce.JobSubmitter: number of splits:1  
18/07/29 02:05:37 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532787160984_0036  
18/07/29 02:05:37 INFO impl.YarnClientImpl: Submitted application application_1532787160984_0036  
18/07/29 02:05:37 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532787160984_00  
18/07/29 02:05:37 INFO mapreduce.Job: Running job: job_1532787160984_0036  
18/07/29 02:05:47 INFO mapreduce.Job: Job job_1532787160984_0036 running in uber mode : false  
18/07/29 02:05:47 INFO mapreduce.Job: map 0% reduce 0%  
18/07/29 02:05:54 INFO mapreduce.Job: map 100% reduce 0%  
18/07/29 02:06:02 INFO mapreduce.Job: map 100% reduce 100%  
18/07/29 02:06:02 INFO mapreduce.Job: Job job_1532787160984_0036 completed successfully  
18/07/29 02:06:02 INFO mapreduce.Job: Counters: 49  
File System Counters  
FILE: Number of bytes read=26  
FILE: Number of bytes written=216079  
FILE: Number of read operations=0  
FILE: Number of large read operations=0
```

```
CPU time spent (ms)=1460  
Physical memory (bytes) snapshot=298491904  
Virtual memory (bytes) snapshot=4118196224  
Total committed heap usage (bytes)=170004480  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=68  
File Output Format Counters  
Bytes Written=6  
[acadgild@localhost ~]$ hadoop fs -cat /songShared output/part-r-00000  
18/07/29 02:06:56 WARN util.NativeCodeLoader: Unable to load native-hadoop  
225 2  
You have new mail in /var/spool/mail/acadgild  
[acadgild@localhost ~]$
```

### **Task 3: 1. Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table**

#### **Creating tables Person**

```
mysql> use simplidb;
Reading table information for completion of table and column names
You can turn off this feature with -A

Database changed
mysql> Create table Person(
  -> person_id INT NOT NULL,
  -> lname varchar(20),
  -> fname varchar(20),
  -> area varchar(20),
  -> city varchar(20),
  -> PRIMARY KEY(person_id));
ERROR 1050 (42S01): Table 'Person' already exists
mysql> describe Person
  -> ;
+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+
| person_id | int(11) | NO | PRI | NULL | |
| lname | varchar(20) | YES | | NULL | |
| fname | varchar(20) | YES | | NULL | |
| area | varchar(20) | YES | | NULL | |
| city | varchar(20) | YES | | NULL | |
+-----+
5 rows in set (0.19 sec)

mysql> select * from Person;
+-----+
| person_id | lname | fname | area | city |
+-----+
| 1 | Shyam | Ram | Patna | Bihar |
| 2 | Tanya | Priya | Whitefield | Bangalore |
| 3 | James | Brown | New York | United States |
| 4 | Jhon | Miller | Los Angeles | United States |
| 789 | a | b | c | d |
+-----+
5 rows in set (0.00 sec)

mysql> update Person set lname ='Shyam' where preson id=1;
ERROR 1054 (42S22): Unknown column 'preson id' in 'where clause'
mysql> update Person set lname ='Shyam' where person_id=1;
Query OK, 0 rows affected (0.00 sec)
Rows matched: 1 Changed: 0 Warnings: 0

mysql> select * from Person;
+-----+
| person_id | lname | fname | area | city |
+-----+
| 1 | Shyam | Ram | Patna | Bihar |
| 2 | Tanya | Priya | Whitefield | Bangalore |
| 3 | James | Brown | New York | United States |
| 4 | Jhon | Miller | Los Angeles | United States |
| 789 | a | b | c | d |
+-----+
5 rows in set (0.00 sec)
```

**Command:** Importing data from table Person into directory /sqoopout

**sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root --password Root@123 target-dir /sqoopout**



```

You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root --password Root@123 --target-dir /sqoopout
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/07/30 07:50:48 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
18/07/30 07:50:48 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
18/07/30 07:50:49 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/07/30 07:50:49 INFO tool.CodeGenTool: Beginning code generation
Mon Jul 30 07:50:49 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements
nection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'.
either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/07/30 07:50:50 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/07/30 07:50:50 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/07/30 07:50:50 INFO orm.CompilationManager: HADOOP MAPRED HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/97fe721b22a2e41ddf8368b9d27031e/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.

```

```

Bytes Written=150
18/07/30 07:51:45 INFO mapreduce.ImportJobBase: Transferred 150 bytes in 48.3787 seconds (3.1005 bytes/sec)
18/07/30 07:51:45 INFO mapreduce.ImportJobBase: Retrieved 5 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -ls /sqoopout
18/07/30 08:48:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 5 items
-rw-r--r-- 1 acadgild supergroup 0 2018-07-30 07:51 /sqoopout/ SUCCESS
-rw-r--r-- 1 acadgild supergroup 138 2018-07-30 07:51 /sqoopout/part-m-00000
-rw-r--r-- 1 acadgild supergroup 0 2018-07-30 07:51 /sqoopout/part-m-00001
-rw-r--r-- 1 acadgild supergroup 0 2018-07-30 07:51 /sqoopout/part-m-00002
-rw-r--r-- 1 acadgild supergroup 12 2018-07-30 07:51 /sqoopout/part-m-00003
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /sqoopout/part-r-00000
18/07/30 08:49:17 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
cat: '/sqoopout/part-r-00000': No such file or directory
[acadgild@localhost ~]$ hadoop fs -ls /sqoopout/part-r-00000
18/07/30 08:49:29 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
ls: '/sqoopout/part-r-00000': No such file or directory
[acadgild@localhost ~]$ hadoop fs -cat /sqoopout/part-m-00000
18/07/30 08:49:51 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Patna,Bihar,Ram,Shyam,1
Whitefiled,Bangalore,Priya,Tanya,2
New York,United States,Brown,James,3
Los Angeles ,United States,Miller,Jhon,4
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /sqoopout/part-m-00003
18/07/30 08:50:01 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
c,d,b,a,789

```

**Comments:** Exporting data present in SQOOPOUT folder into database

**Command:** `sqoop export --connect jdbc:mysql://localhost/simplidb --username 'root' -P --table 'Person' --export-dir '/sqoopout' --input-fields-terminated-by ',' -m 1 --columns person_id,lname,fname,area,city`



```

[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/simplidb --username 'root' -P --table 'Person' --export-dir '/sqoopout'
city,area,fname,lname,person_id
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO HOME to the root of your Accumulo installation.
18/07/30 12:26:11 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/07/30 12:26:23 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/07/30 12:26:23 INFO tool.CodeGenTool: Beginning code generation
Mon Jul 30 12:26:24 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5 security guidelines, secure connections to MySQL servers should be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate option is disabled by default.

```

```

Total time spent by all map tasks (ms)=7350
Total vcore-milliseconds taken by all map tasks=7350
Total megabyte-milliseconds taken by all map tasks=7526400
Map-Reduce Framework
  Map input records=4
  Map output records=4
  Input split bytes=304
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=76
  CPU time spent (ms)=1630
  Physical memory (bytes) snapshot=105508864
  Virtual memory (bytes) snapshot=2061332480
  Total committed heap usage (bytes)=32571392
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/07/30 12:26:56 INFO mapreduce.ExportJobBase: Transferred 451 bytes in 25.5733 seconds (17.6356 bytes/sec)
18/07/30 12:26:56 INFO mapreduce.ExportJobBase: Exported 4 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

```

mysql> delete from Person where person_id =4;
Query OK, 1 row affected (0.00 sec)

mysql> select * from Person;
Empty set (0.00 sec)

mysql> select * from Person;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area | city |
+-----+-----+-----+-----+-----+
| 1 | Shyam | Ram | Bihar | Patna |
| 2 | Tanya | Priya | Bangalore | Whitefield |
| 3 | James | Brown | United States | Newyork |
| 4 | John | Miller | United States | Los Angeles |
+-----+-----+-----+-----+-----+
4 rows in set (0.00 sec)

mysql>

```

## 2. Use Sqoop tool to import data from the same mysql table where the person\_id = 3 into a new hdfs directory SQOOP\_FILTER

**Comments:** Executing command to import data with person\_id="3" from table "Person" into directory "/SQOOP\_FILTER" using below command.

**sqoop import --connect jdbc:mysql://localhost/simplidb --table 'Person' --username root -P --target-dir '/SQOOP\_FILTER' -m 1 --where person\_id=3**

```
You have new mail in /var/spool/mail/acadgild
acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table 'Person' --username root -P --target-dir '/SQOOP_FILTER' -m 1 --where person_id=3
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/07/30 12:35:03 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/07/30 12:35:09 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/07/30 12:35:09 INFO tool.CodeGenTool: Beginning code generation
Mon Jul 30 12:35:09 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements, a secure connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You can either explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/07/30 12:35:10 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/07/30 12:35:10 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/07/30 12:35:10 INFO mapreduce.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
```

```
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=125
CPU time spent (ms)=2240
Physical memory (bytes) snapshot=107327488
Virtual memory (bytes) snapshot=2063437824
Total committed heap usage (bytes)=32571392

File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=36
18/07/30 12:35:44 INFO mapreduce.ImportJobBase: Transferred 36 bytes in 29.0579 seconds (1.24 B/sec)
18/07/30 12:35:44 INFO mapreduce.ImportJobBase: Retrieved 1 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -ls /SQOOP_FILTER
18/07/30 12:36:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java-based ones
Found 2 items
-rw-r--r--  1 acadgild supergroup          0 2018-07-30 12:35 /SQOOP_FILTER/_SUCCESS
-rw-r--r--  1 acadgild supergroup        36 2018-07-30 12:35 /SQOOP_FILTER/part-m-00000
[acadgild@localhost ~]$ hadoop fs -cat /SQOOP_FILTER/part-m-00000
18/07/30 12:36:32 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java-based ones
United States,Newyork,Brown,James,3
[acadgild@localhost ~]$
```